

Attention is all you need!



Sommaire

- Introduction
- Détails du modèle
 - Encoder
 - Decoder
- Résultats
- Conclusion



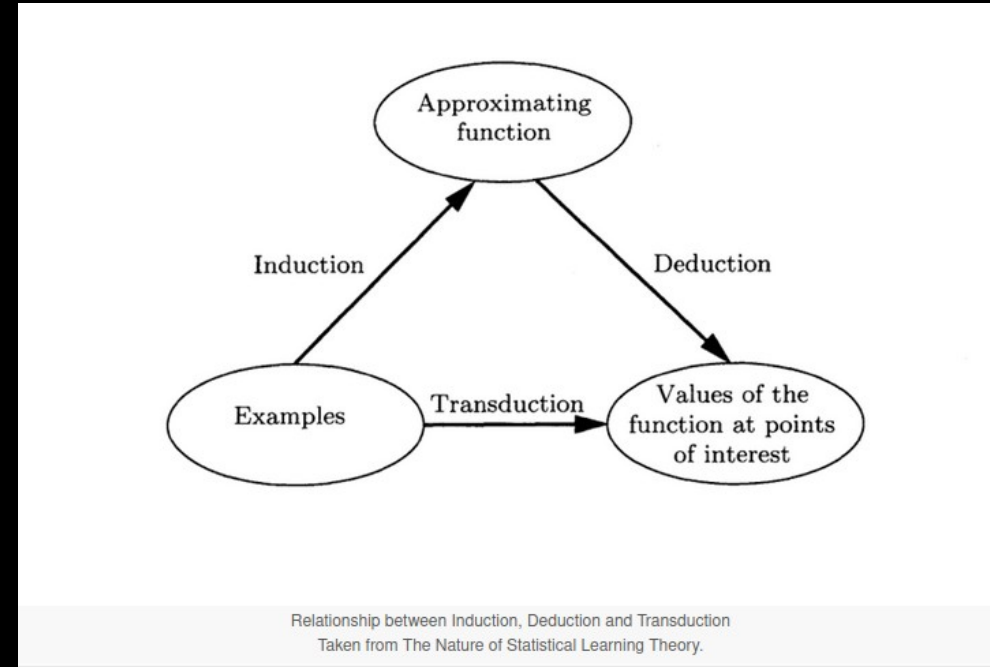
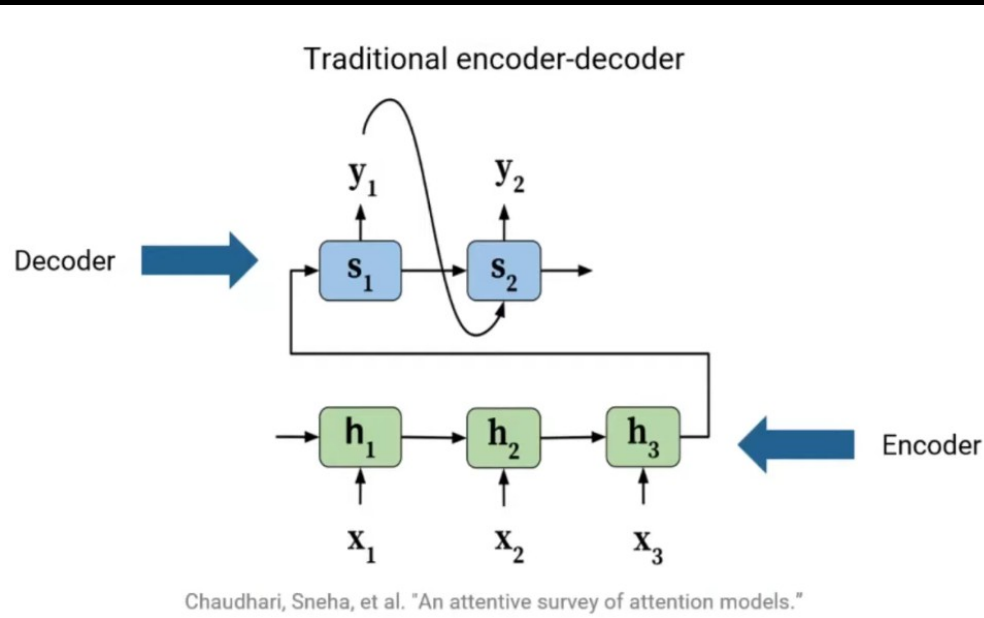
Introduction



Introduction


Transduction et modélisation de séquences :

- Modélisation du langage
- Traduction
- RNNs, LSTM, gated RNNs



Introduction

Encoder-decoder model with Attention



Chaudhari, Sneha, et al. "An attentive survey of attention models."

- 2014 : Neural Machine Translation by Jointly Learning to Align and Translate" by Bahdanau et al.
- Affinités entre éléments distants
- Mais RNNs trop lents (GPU...)



Introduction



Cornell University

arXiv > cs > arXiv:1706.03762

Computer Science > Computation and Language

[Submitted on 12 Jun 2017 (v1), last revised 6 Dec 2017 (this version, v5)]

Attention Is All You Need

Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, Illia Polosukhin

2017 : Introduction d'une nouvelle architecture : les transformers.
Pas de traitement séquentiel --> //



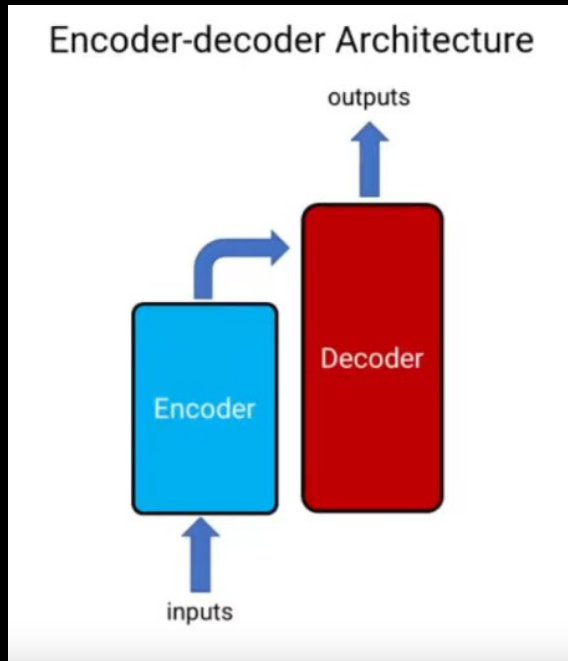
Détails du modèle Encoder & Decoder



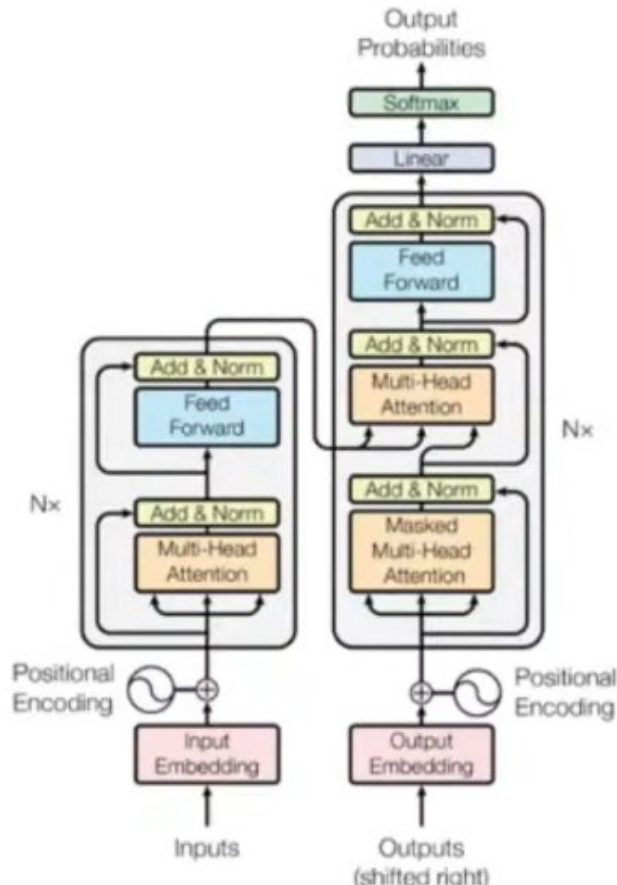
Détails de l'encoder Preprocessing...



Détails de l'encoder : preprocessing

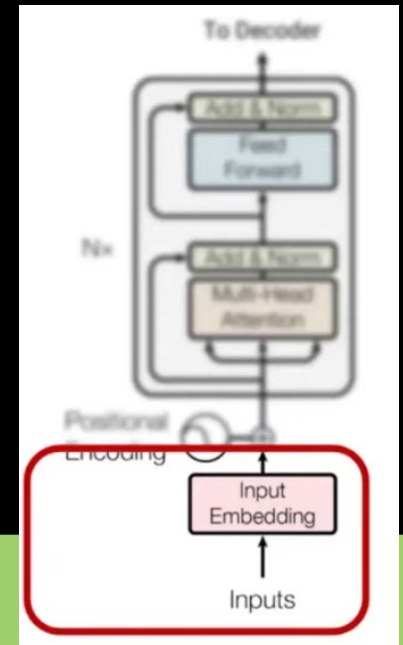


Transformer Architecture

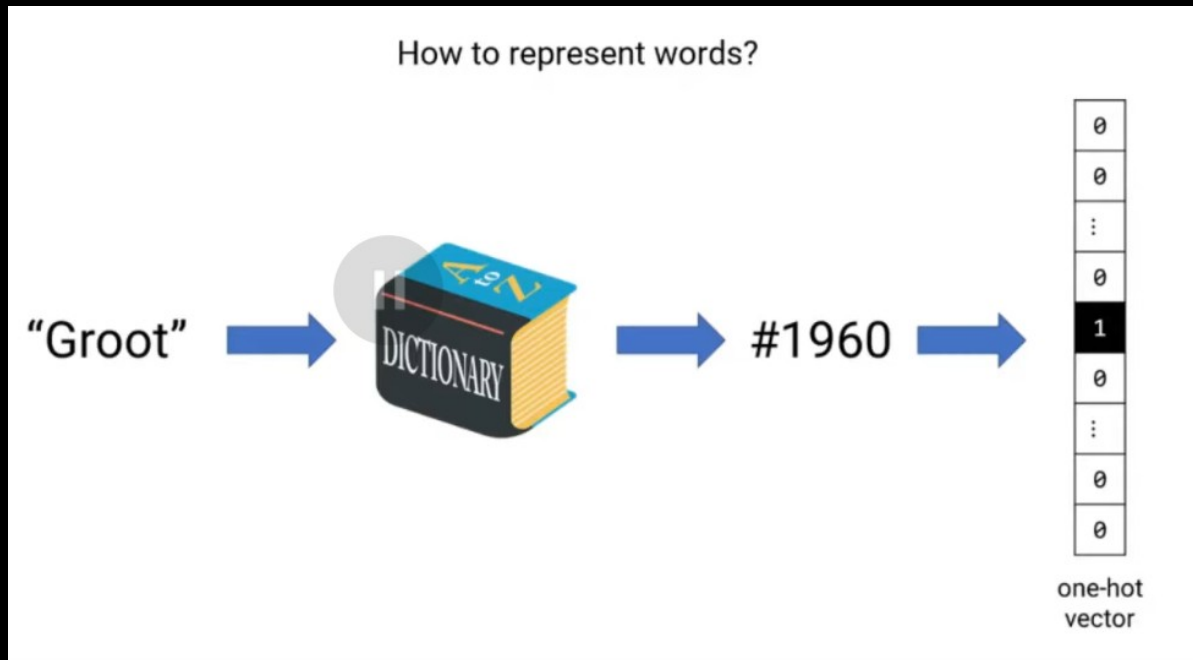


Inputs = séquence de mots

Comment encoder ces mots?

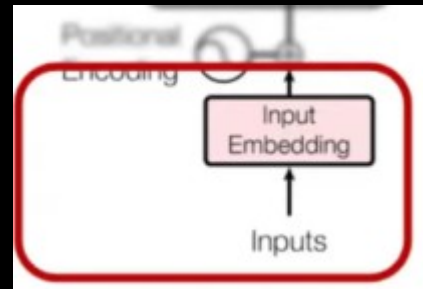


Détails de l'encoder : preprocessing



Première idée : one-hot vector

-- > One hot encoder



One-Hot Encoding

datagy.io

Island	Biscoe	Dream	Torgensen
Biscoe	1	0	0
Torgensen	0	0	1
Dream	0	1	0



Détails de l'encoder : preprocessing

```
# Loading a Sample DataFrame
import pandas as pd

df = pd.DataFrame({
    'Name': ['Joan', 'Matt', 'Jeff', 'Melissa', 'Devi'],
    'Gender': ['Female', 'Male', 'Male', 'Female', 'Female'],
    'House Type': ['Apartment', 'Detached', 'Apartment', None,
'Semi-Detached']
})
```

```
print(df)
```

```
# Returns:
```

#	Name	Gender	House Type
# 0	Joan	Female	Apartment
# 1	Matt	Male	Detached
# 2	Jeff	Male	Apartment
# 3	Melissa	Female	None
# 4	Devi	Female	Semi-Detached

```
# Understanding the Pandas get_dummies
```

```
import pandas as pd
```

```
pd.get_dummies(
```

```
    data,
```

```
    prefix=None,
```

```
    prefix_sep='_',
```

```
    dummy_na=False,
```

```
    columns=None,
```

`data=` represents the data from which to get the dummy indicators (either array-like, Pandas Series, or Pandas DataFrame)

`prefix=` represents the string to append to DataFrame column names

`prefix_sep=` represents what delimiter to use

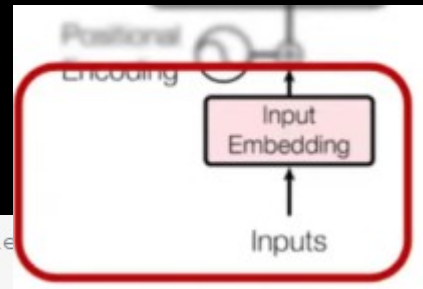
`dummy_na=` represents whether to add a column or not for missing values

`columns=` represents the names of the columns to be encoded

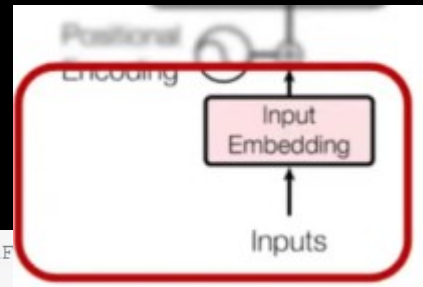
`sparse=` represents whether the data should be a sparse array or a regular [NumPy](#) array

`drop_first=` represents whether to drop the first level or not

`dtype=` represents the data type for new columns



Détails de l'encoder : preprocessing



```
# One-Hot Encoding a Single DataFrame Series
```

```
import pandas as pd
```

```
df = pd.DataFrame({
    'Name': ['Joan', 'Matt', 'Jeff', 'Melissa', 'Devi'],
    'Gender': ['Female', 'Male', 'Male', 'Female', 'Female'],
    'House Type': ['Apartment', 'Detached', 'Apartment', None,
'Semi-Detached']
})
```

```
print(pd.get_dummies(df['Gender']))
```

```
# Returns:
```

#	Female	Male
# 0	1	0
# 1	0	1
# 2	0	1
# 3	1	0
# 4	1	0

```
# One-Hot Encoding and Returning a DataFrame
```

```
import pandas as pd
```

```
df = pd.DataFrame({
    'Name': ['Joan', 'Matt', 'Jeff', 'Melissa', 'Devi'],
    'Gender': ['Female', 'Male', 'Male', 'Female', 'Female'],
    'House Type': ['Apartment', 'Detached', 'Apartment', None,
'Semi-Detached']
})
```

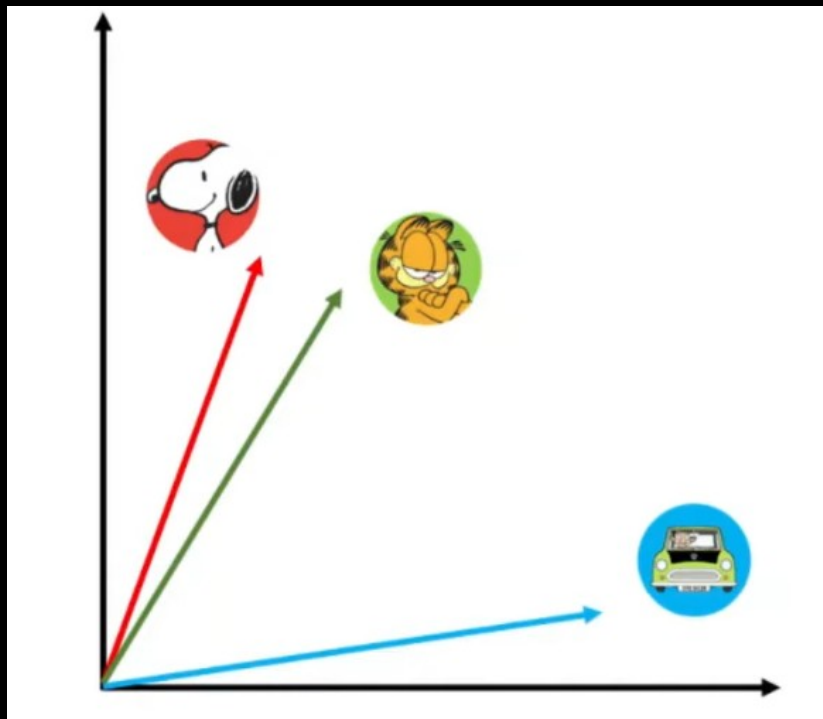
```
ohe = pd.get_dummies(data=df, columns=['Gender'])
```

```
print(ohe)
```

```
# Returns:
```

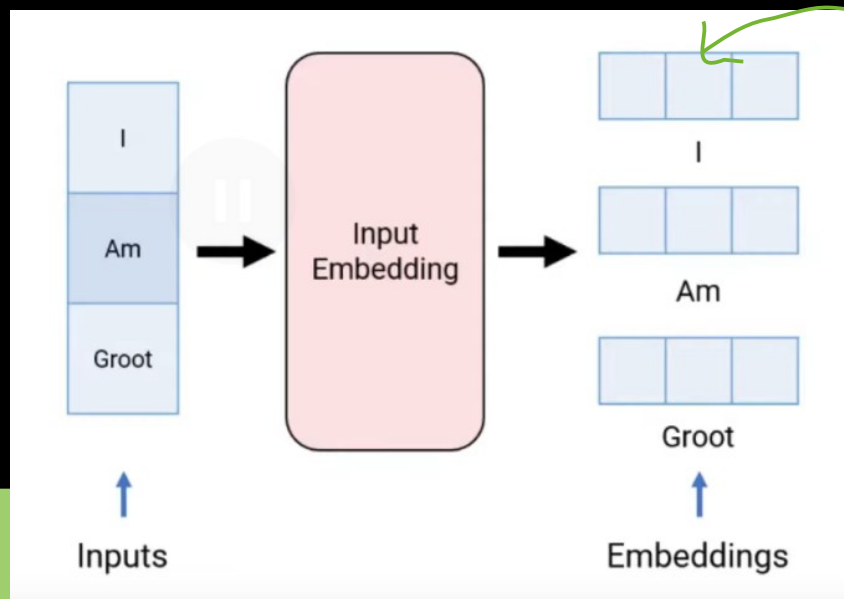
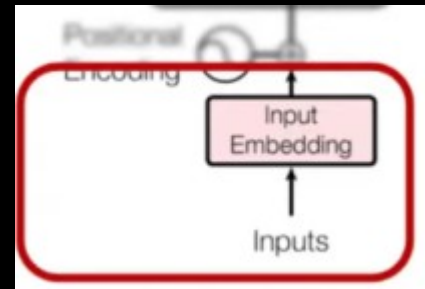
#	Name	House Type	Gender_Female	Gender_Male
# 0	Joan	Apartment	1	0
# 1	Matt	Detached	0	1
# 2	Jeff	Apartment	0	1
# 3	Melissa	None	1	0
# 4	Devi	Semi-Detached	1	0

Détails de l'encoder : preprocessing



Meilleure idée : word embeddings

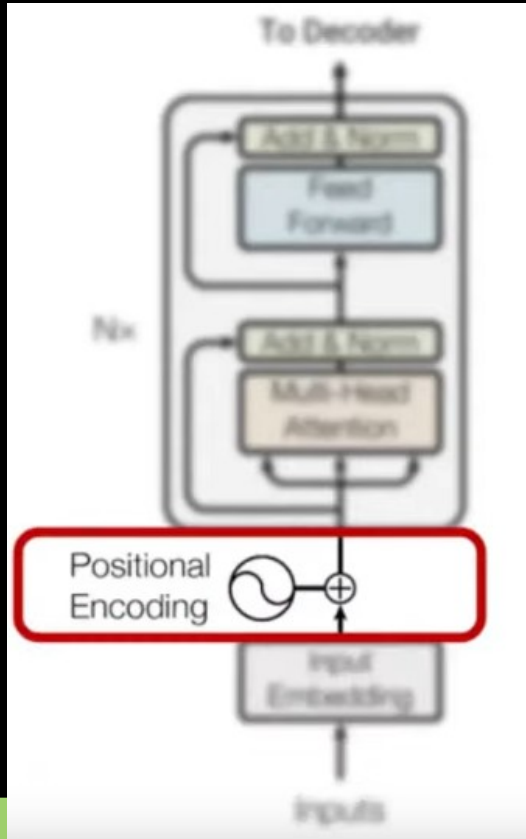
Mots similaires \leftrightarrow
représentations similaires



512



Détails de l'encoder : preprocessing



Positional encoding = rajouter de l'info sur les positions relatives des mots

-- > Faire la somme d'un nouveau vecteur (positions relatives) et des vecteurs précédents (mots)

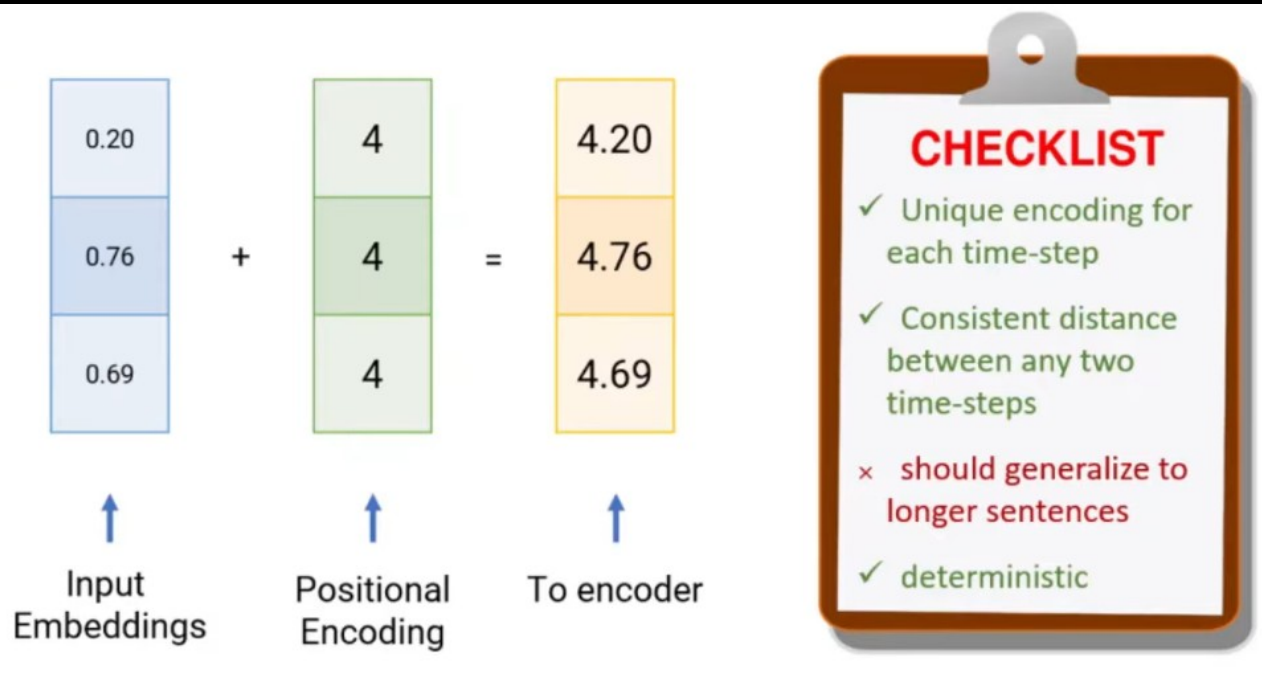
CHECKLIST

- Unique encoding for each time-step
- Consistent distance between any two time-steps
- should generalize to longer sentences
- deterministic

The Avengers defeated Thanos

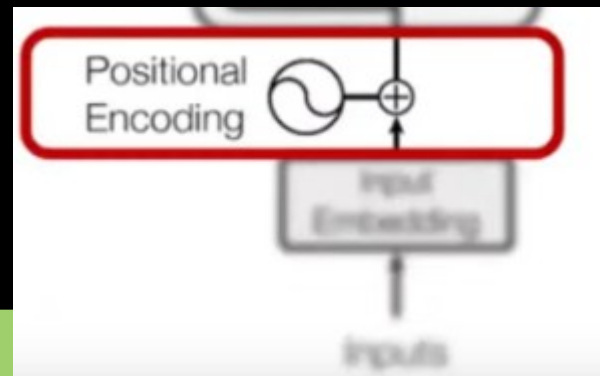
Thanos defeated The Avengers

Détails de l'encoder : preprocessing

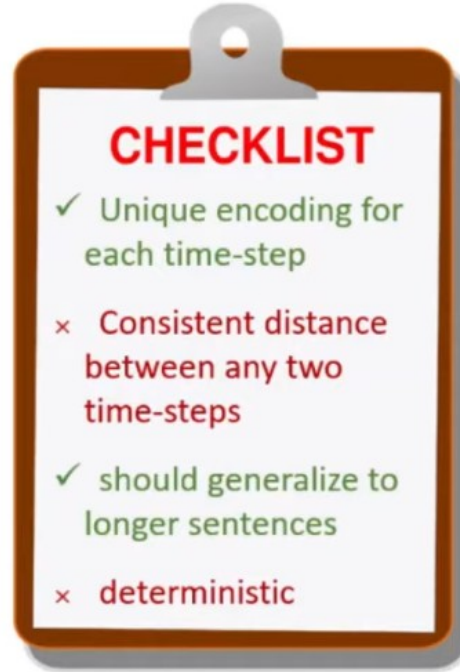
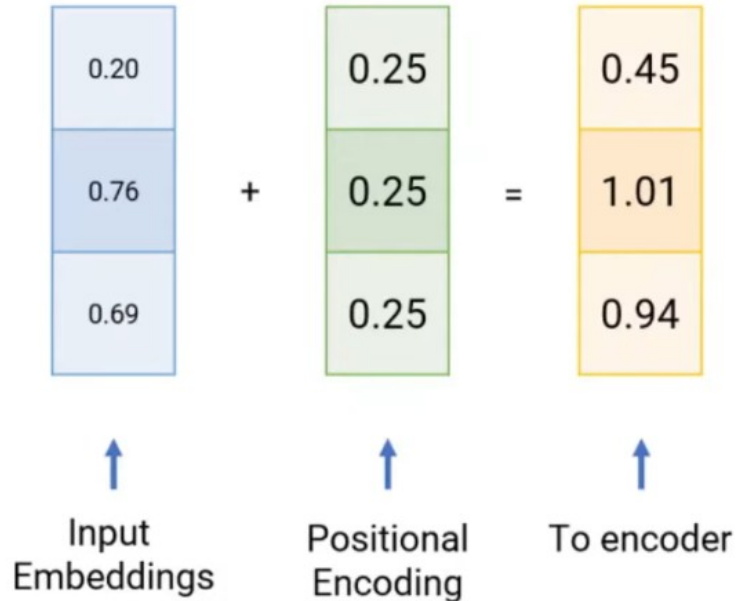


Rajouter 1 à la somme des vecteurs précédents pour chaque nouveau mot...

>>

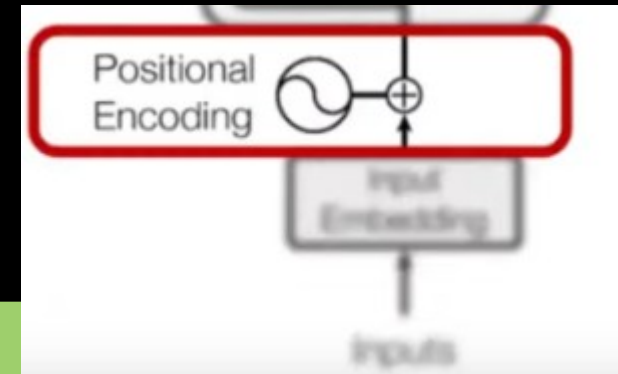


Détails de l'encoder : preprocessing



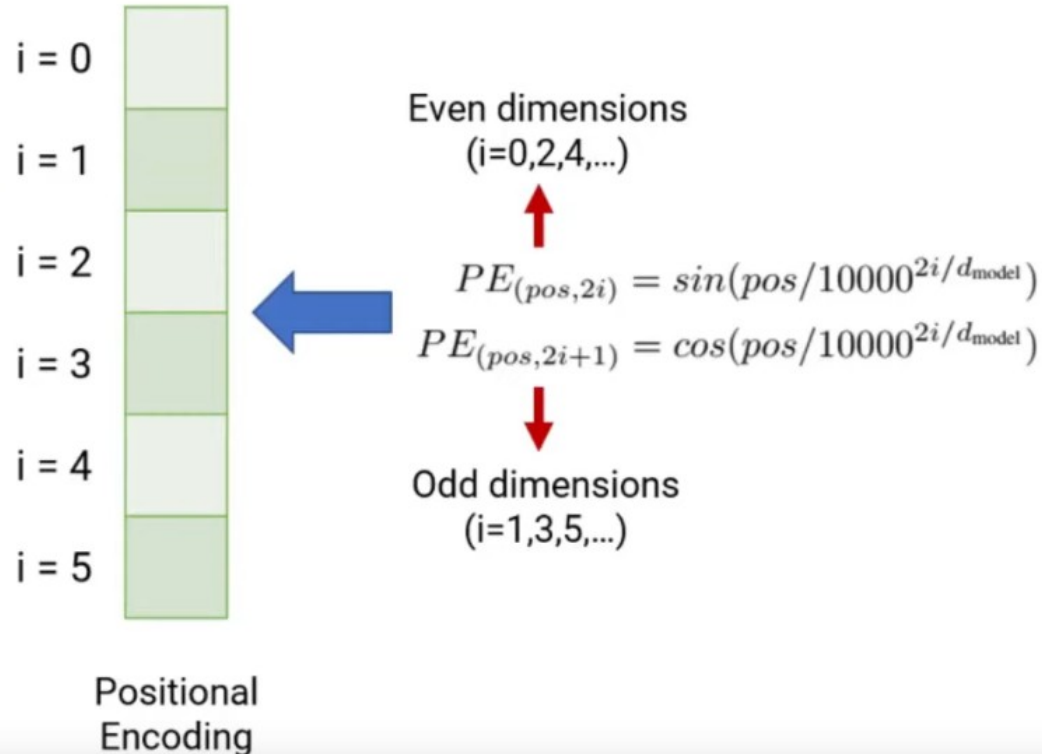
Rajouter un nombre entre 0 et 1

!deterministe <- plusieurs mots / delta T



Détails de l'encoder : preprocessing


How to create the Positional Encoding (Transformer style)



Vecteur de 512 lignes
 $d_{model}=512$

$$\vec{p}_t = \begin{bmatrix} \sin(\omega_1 \cdot t) \\ \cos(\omega_1 \cdot t) \\ \sin(\omega_2 \cdot t) \\ \cos(\omega_2 \cdot t) \\ \vdots \\ \sin(\omega_{d/2} \cdot t) \\ \cos(\omega_{d/2} \cdot t) \end{bmatrix}_{d \times 1} \quad \left(\omega_k = \frac{1}{10000^{\frac{2k}{d}}} \right)$$

https://kazemnejad.com/blog/transformer_architecture_positional_encoding/



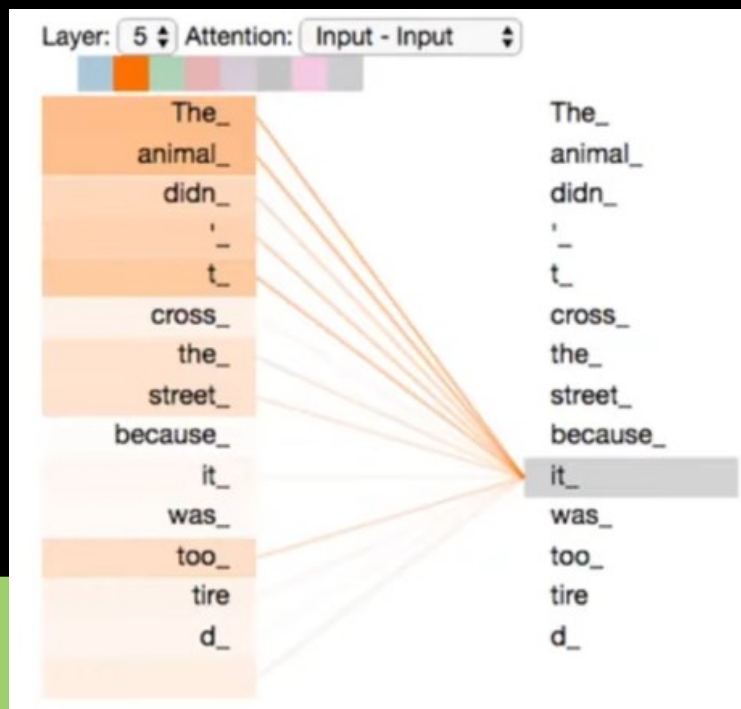

Fin du preprocessing de l'encoder!



Détails de l'encoder : Multi-head self-attention mechanism



Détails de l'encoder : self attention



Détails de l'encoder : self attention

